



北京邮电大学软件学院

School Of software Engineering Of BUPT

厚德 博学 敬业 乐群



Operating Systems

Lecture 3 Threads

Jinpengchen

Email: jpchen@bupt.edu.cn



Catalog Description

- ⊕ Overview
- ⊕ Multithreading Models
- ⊕ Threading Issues



Why Processes?

- ✚ 自从60年代提出进程概念以来，在操作系统中一直都是以进程作为独立运行的基本单位，直到80年代中期，人们又提出了更小的能独立运行的基本单位—线程。



Why Processes?

- ✚ [案例]编写一个MP3播放软件，核心功能模块有三个：
 - ✚ 从MP3音频文件中读取数据；
 - ✚ 对数据进行解压缩；
 - ✚ 把解压缩后的音频数据播放出来。



Why Processes?

单进程的实现方法:

```
I/O  Main()
CPU  {
      While(true)
      {read();
       decompress();
       Play();
      }
    }
    read(){...};
    decompress(){...};
    Play(){...};
```

问题:

- ✓ 播放出来的声音是否连贯?
- ✓ 各个函数之间不是并发执行, 影响资源的使用效率?

Youtube的状态栏



Why Processes?

多进程的实现方法:

程序1

```
Main()
```

```
{
```

```
While(true)
```

```
{read();
```

```
}
```

```
}
```

```
read(){...};
```

程序2

```
Main()
```

```
{
```

```
While(true)
```

```
{decompress();
```

```
}
```

```
}
```

```
decompress(){...};
```

程序3

```
Main()
```

```
{
```

```
While(true)
```

```
{Play();
```

```
}
```

```
}
```

```
Play(){...};
```



Why Processes?

✚ 怎么解决这些问题:

需要提出一种新的实体，满足以下特性：

实体之间可以并发地执行；

实体之间共享相同的地址空间；

这种实体就是：线程（Thread）



Thread concept overview

✿ What is the thread?

❏ Thread:

- ✓ A sequential execution stream within a process
- ✓ a thread of execution
- ✓ 进程当中的一条执行流程



Thread concept overview


- ❁ A thread is a basic unit of CPU utilization;
 - ❁ it comprises a thread ID, a program counter, a register set, and a stack.
 - ❁ it shares with other threads belonging to the same process, the code section, the data section, and other OS resources, such as open files, signals, etc
- ❁ A traditional process has a single thread of control: heavyweight process.



Thread concept overview

从两个方面来理解进程

- 从资源组合的角度—进程把一组相关的资源组合起来，构成一个资源平台（环境），包括地址空间（代码段、数据段）、打开的文件等各种资源
- 从运行的角度：代码在这个平台上的一条执行流程（线程）






Thread concept overview

❁ 进程 = 线程 + 资源平台

❁ 优点:

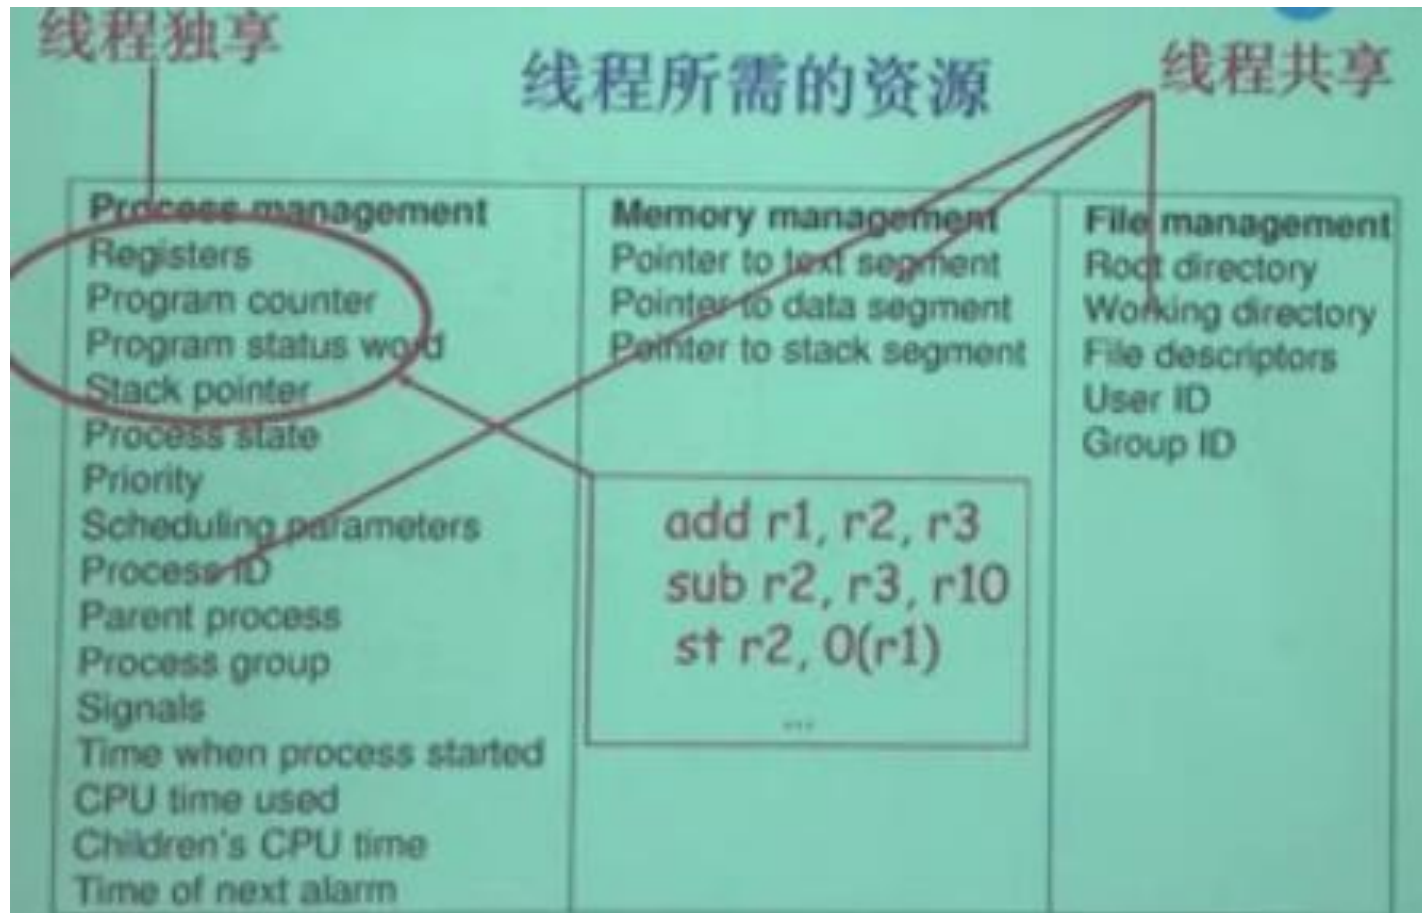
- ✓ 一个进程中可以同时存在多个线程
- ✓ 各个线程之间可以并发地执行
- ✓ 各个线程之间可以共享地址空间





Thread concept overview

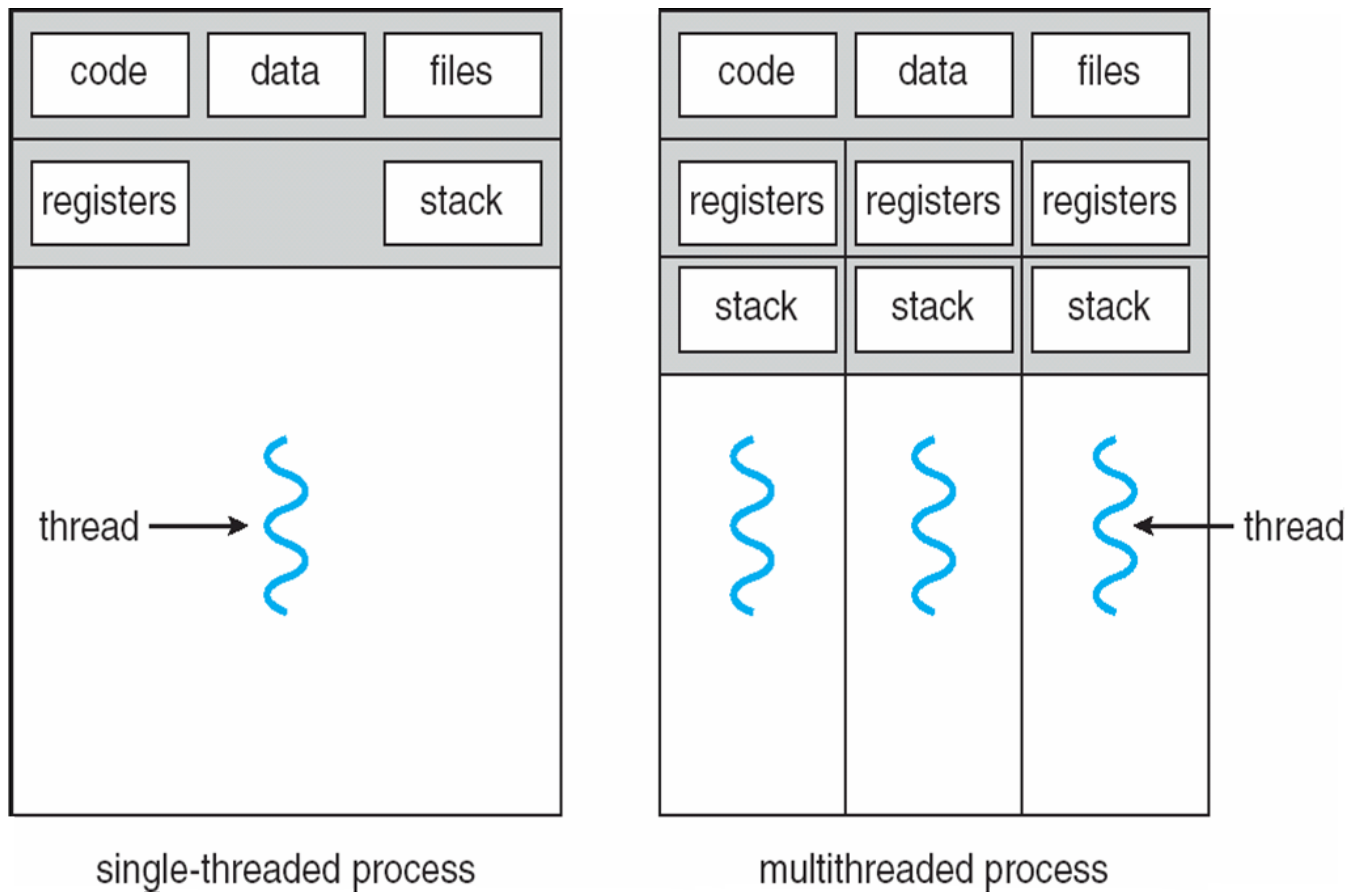
Single and Multithreaded Processes





Thread concept overview

Single and Multithreaded Processes





Thread concept overview

❁ 线程与进程的比较

- ❁ 进程是资源分配单位, 线程是CPU分配单位
- ❁ 进程拥有一个完整的资源平台, 而线程值独享必不可少的资源, 如寄存器和栈
- ❁ 线程同样具有就绪、阻塞和执行三种状态, 同样具有状态之间的转换关系
- ❁ 线程能减少并发执行的时间和空间开销
- ❁ 线程=轻量级进程 (lightweight process)



Thread concept overview

✚ Motivation

- ✚ On modern desktop PC, many APPs are multithreaded.
- ✚ a separate process with several threads
- ✚ Example 1: A web browser
 - ✓ one for displaying images or text;
 - ✓ another for retrieving data from network
- ✚ Example 2: A word processor
 - ✓ one for displaying graphics;
 - ✓ another for responding to keystrokes from the user;
 - ✓ and a third for performing spelling & grammar checking in the background



Thread concept overview

✚ Motivation

- ✚ In certain situations, a single application may be required to perform several similar tasks. Example: a web server
- ✚ Allow a server to service several concurrent requests. Example: an RPC server and Java's RMI systems
- ✚ The OS itself needs to perform some specific tasks in kernel, such as managing devices or interrupt handling.
 - ✓ PARTICULAR, many OS systems are now multithreaded.
 - ✓ Example: Solaris, Linux



Thread concept overview

✚ Benefits

✚ Responsiveness (响应度高)

✓ Example: an interactive application such as web browser, while one thread loading an image, another thread allowing user interaction

✚ Resource Sharing

✓ address space, memory, and other resources

✚ Economy

✓ Solaris:

✓ creating a process is about 30 times slower than creating a thread;

✓ context switching is about 5 times slower

✚ Utilization of MP Architectures

✓ parallelism and concurrency ↑



Catalog Description

- ⊕ Overview
- ⊕ Multithreading Models
- ⊕ Threading Issues



Two methods to support threads

☀ User threads VS. Kernel threads

⊞ User threads

- ✓ Thread management done by user-level threads library without kernel support

- Kernel may be multithreaded or not.

⊞ Three primary thread libraries:

- ✓ POSIX Pthreads

- ✓ Win32 threads

- ✓ Java threads



Two methods to support threads

🌀 User threads VS. Kernel threads

⌘ Kernel threads

- ✓ Supported by the Kernel, usually may be slower than user thread

⌘ Examples:

- ✓ Windows XP/2000
- ✓ Solaris
- ✓ Linux
- ✓ Tru64 UNIX (formerly Digital UNIX)
- ✓ Mac OS X



Multithreading Models

✚ The relationship between user threads and kernel threads

- ✚ Many-to-One [n:1]

- ✚ One-to-One [1:1]

- ✚ Many-to-Many [n:m]


✚ Many-to-One [n:1]

- ✚ Many user-level threads
Mapped to single kernel thread

- ✚ Examples:

- ✓ Solaris Green Threads

- ✓ GNU Portable Threads






Multithreading Models

One-to-One [1:1]

- Each user-level thread maps to a kernel thread
- Examples:
 - ✓ Windows NT/XP/2000
 - ✓ Linux
 - ✓ Solaris 9 and later






Multithreading Models

❖ Many-to-Many [n:m]

- ❑ Allows many user level threads to be mapped to many kernel threads
- ❑ Allows the operating system to create a sufficient number of kernel threads
- ❑ Examples:
 - ✓ Solaris prior to version 9
 - ✓ Windows NT/2000 with the ThreadFiber package






Multithreading Models

✦ Two-level Model, a popular variation on many-to-many model

✦ Similar to $n:m$, except that it allows a user thread to be bound to a kernel thread

✦ Examples:

- ✓ IRIX
- ✓ HP-UX
- ✓ Tru64 UNIX
- ✓ Solaris 8 and earlier





Catalog Description

- ⊕ Overview
- ⊕ Multithreading Models
- ⊕ Threading Issues



Threading Issues

- ✚ Semantics of `fork()` and `exec()` system calls
 - ✚ Does `fork()` duplicate only the calling thread or all threads?
 - ✚ Some UNIX system have chosen to have two versions
 - ✚ Which one version to use? Depend on the APP.
- ✚ Thread cancellation
 - ✚ Terminating a thread before it has finished
 - ✚ Two general approaches:
 - ✓ Asynchronous (异步) cancellation terminates the target thread immediately
 - ✓ Deferred (延时) cancellation allows the target thread to periodically check if it should be cancelled



Threading Issues

✦ Signal Handling

- ✦ Signals are used in UNIX systems to notify a process that a particular event has occurred:
 - ✓ Synchronous: illegal memory access, division by 0
 - ✓ Asynchronous: Ctrl+C
- ✦ All signals follow the same pattern:
 - ✓ Signal is generated by particular event
 - ✓ Signal is delivered to a process
 - ✓ Signal is handled
- ✦ Signal handler may be handled by
 - ✓ a default signal handler, or
 - ✓ a user-defined signal handler



Threading Issues

✦ Signal Handling

- ✦ When multithread, where should a signal be delivered?
 - ✓ Deliver the signal to the thread to which the signal applies
 - ✓ Deliver the signal to every thread in the process
 - ✓ Deliver the signal to certain threads in the process
 - ✓ Assign a specific thread to receive all signals for the process



Threading Issues

✚ Thread Pools

- ✚ Create a number of threads in a pool where they await work
- ✚ Advantages:
 - ✓ Usually slightly faster to service a request with an existing thread than create a new thread
 - ✓ Allows the number of threads in the application(s) to be bound to the size of the pool

✚ Thread Specific Data

- ✚ Allows each thread to have its own copy of data
- ✚ Useful when you do not have control over the thread creation process (i.e., when using a thread pool)



Threading Issues

✦ Scheduler Activations

- ✦ Both n:m and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- ✦ Scheduler activations provide up calls – a communication mechanism from the kernel to the thread library
- ✦ This communication allows an application to maintain the correct number kernel threads