



北京邮电大学软件学院

School Of software Engineering Of BUPT

厚德 博学 敬业 乐群



# *Operating Systems*

---

## **Lecture 12 : IO Systems**

**Jinpengchen**

**Email: [jpchen@bupt.edu.cn](mailto:jpchen@bupt.edu.cn)**



# *Catalog Description*

---

- ✿ I/O Hardware
- ✿ Application I/O Interface
- ✿ Kernel I/O Subsystem
- ✿ Transforming I/O Requests to Hardware Operations
- ✿ Performance



# Overview

---

- ✿ I/O devices
  - ✦ vary widely
- ✿ The control of devices connected to the computer is a major concern of OS designers.
- ✿ How OS manages and controls various peripherals(外设)?



# *I/O Hardware*

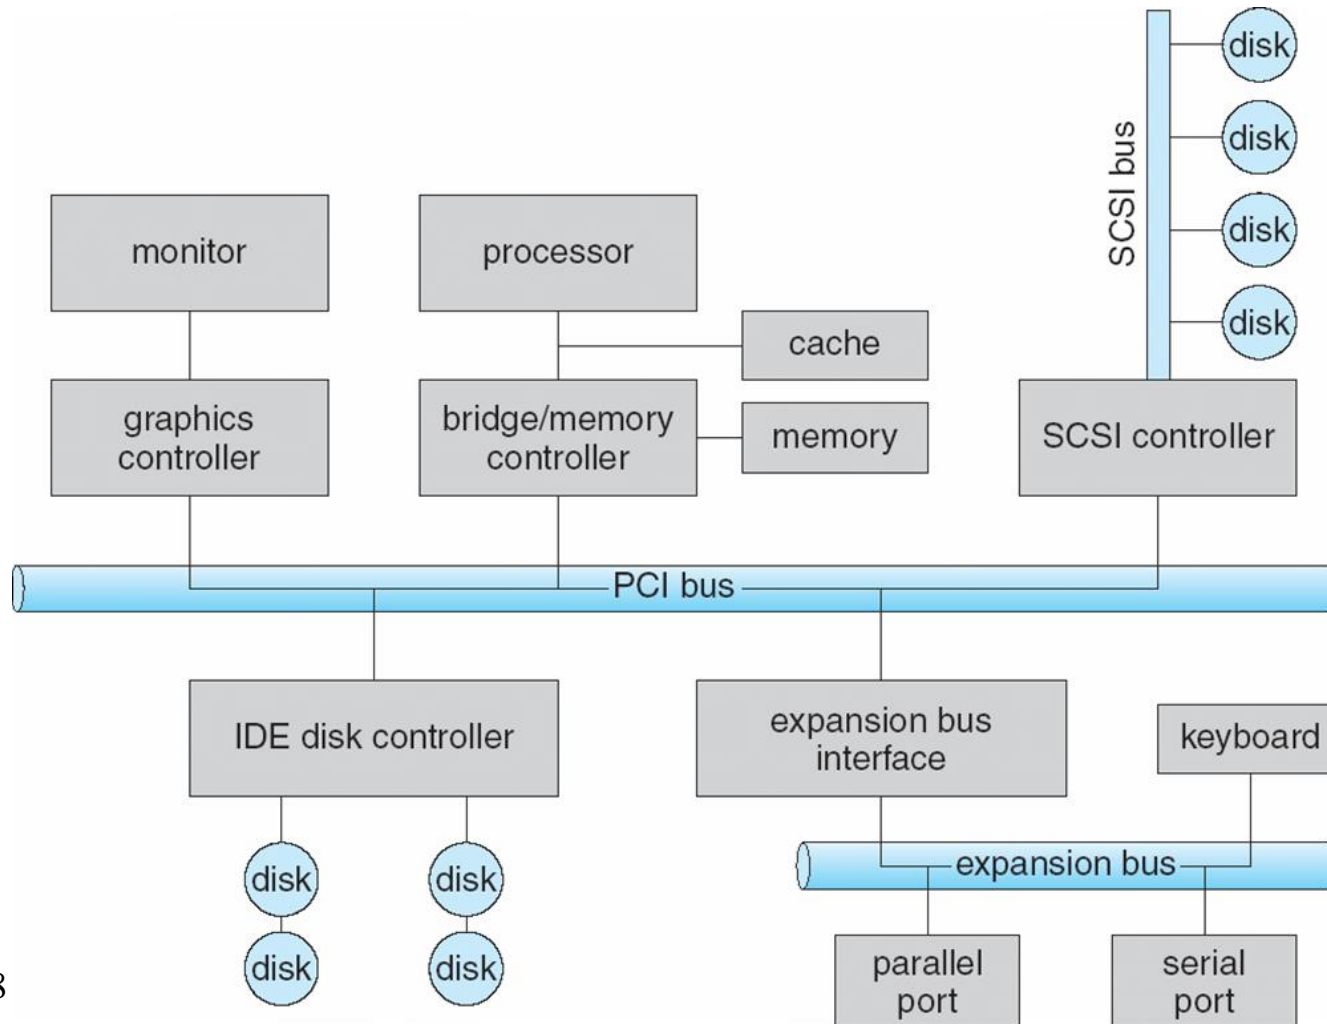
---

- ✿ I/O Hardware
  - ✦ Polling (轮询方式)
  - ✦ Interrupts (中断方式)
  - ✦ Direct Memory Access (DMA方式)
  - ✦ I/O hardware summary



# I/O Hardware

## ❁ Incredible variety of I/O devices





# *I/O Hardware*

- ❁ Common concepts : CPU→PORT→BUS→Controller
  - ❁ Port (端口)
  - ❁ Bus (总线) (daisy chain(菊花链) or shared direct access)
    - ✓ PCI (Peripheral Component Interconnect(外部器件互连) )
    - ✓ SCSI (Small computer systems interface)
    - ✓ Expansion bus
  - ❁ Controller (控制器) (host adapter)
- ❁ How can the processor command controller?
  - ❁ Controller has one or more registers for data and control signals.
  - ❁ The processor communicates with the controller by reading and writing bit patterns in the registers.



# *I/O Hardware*

---

- ❁ Two communication techniques:
  - ❁ Direct I/O instructions
    - ✓ Access the port address
    - ✓ Each port typically contains of four registers, i.e., status, control, data-in and data-out.
    - ✓ Instructions: In, out
  - ❁ Memory-mapped I/O
    - ✓ Example: 0xa0000 ~ 0xfffff are reserved to ISA graphics cards and BIOS routines
- ❁ Some systems use both techniques.



# I/O Hardware

## ✚ I/O address range

### ✚ Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)





# *Polling* (轮询方式)

- ✿ Need handshaking (握手)
- ✿ State of device
  - ✦ **command-ready**
    - ✓ In command register
    - ✓ 1: a command is available for the controller
  - ✦ **Busy**
    - ✓ In status register
    - ✓ 0: ready for the next command; 1: busy
  - ✦ **Error**
    - ✓ To indicate whether an I/O is ok



# *Polling* (轮询方式)

- ✿ Basic handshaking notion for writing output
  - ✦ Host repeatedly reads the busy bit until it is 0
  - ✦ Host sets write bit in command register and writes a byte into data-out register
  - ✦ Host sets command-ready bit
  - ✦ When controller notices command-ready, sets busy bit
  - ✦ Controller gets write command and data, and works
  - ✦ Controller clears command-ready bit, error bit and busy bit
- ✿ Step1: Busy-wait cycle to wait for I/O from device  
≡polling



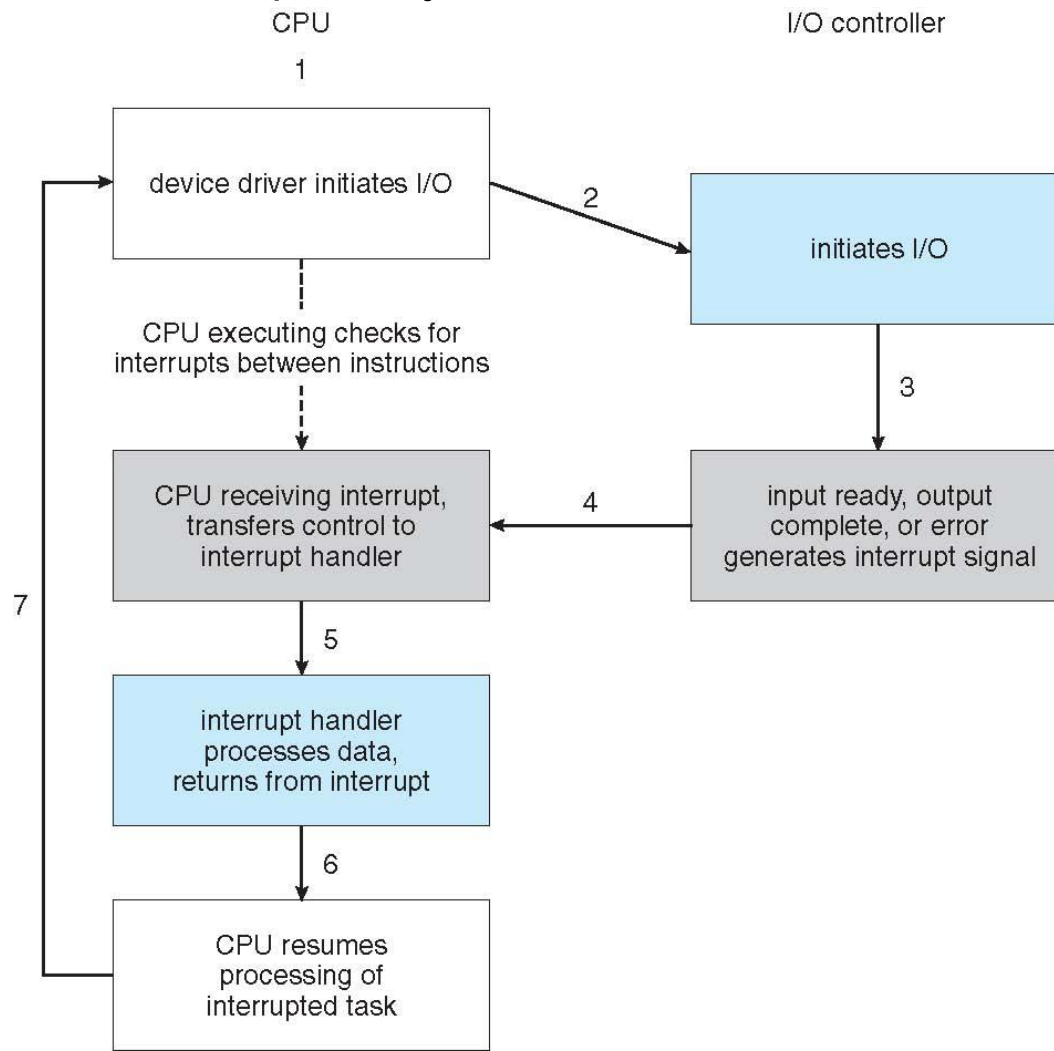
# *Interrupts (中断方式)*

- ⊙ CPU Interrupt-request line triggered by I/O device
- ⊙ Interrupt handler receives interrupts
- ⊙ Basic interrupt scheme
  - ⊞ Raise → Catch → Dispatch → Clear



# *Interrupts (中断方式)*

## Interrupt-Driven I/O Cycle





# *Interrupts (中断方式)*

---

- ✿ More sophisticated interrupt-handling features:
- ✿ Most CPU have two interrupt request line.
  - ✦ Nonmaskable
  - ✦ Maskable to ignore or delay some interrupts
- ✿ Efficient dispatching without polling the devices
  - ✦ Interrupt vector: to dispatch interrupt to correct handler
  - ✦ Interrupt chaining: to allow more device & more interrupt handlers
- ✿ Distinguish between high- and low-priority interrupts:
  - ✦ Interrupt priority: the handling of low-priority interrupts is deferred without masking, even preempted.
- ✿ Interrupt mechanism also used for exceptions



# *Direct Memory Access (DMA方式)*

## ⊗ Direct Memory Access (DMA方式):

Used to avoid programmed I/O for **large data movement**, and bypasses CPU to **transfer data directly between I/O device and memory**

## ⊗ Requires **DMA controller**

- ⊠ the host prepares a **DMA command block** in memory
  - ✓ a pointer to the source of a transfer
  - ✓ a pointer to the destination of the transfer
  - ✓ a count of the number of bytes to be transferred
- ⊠ CPU writes the address of the DMA command block to DMA controller, and then goes on with other work.



# *Direct Memory Access (DMA方式)*

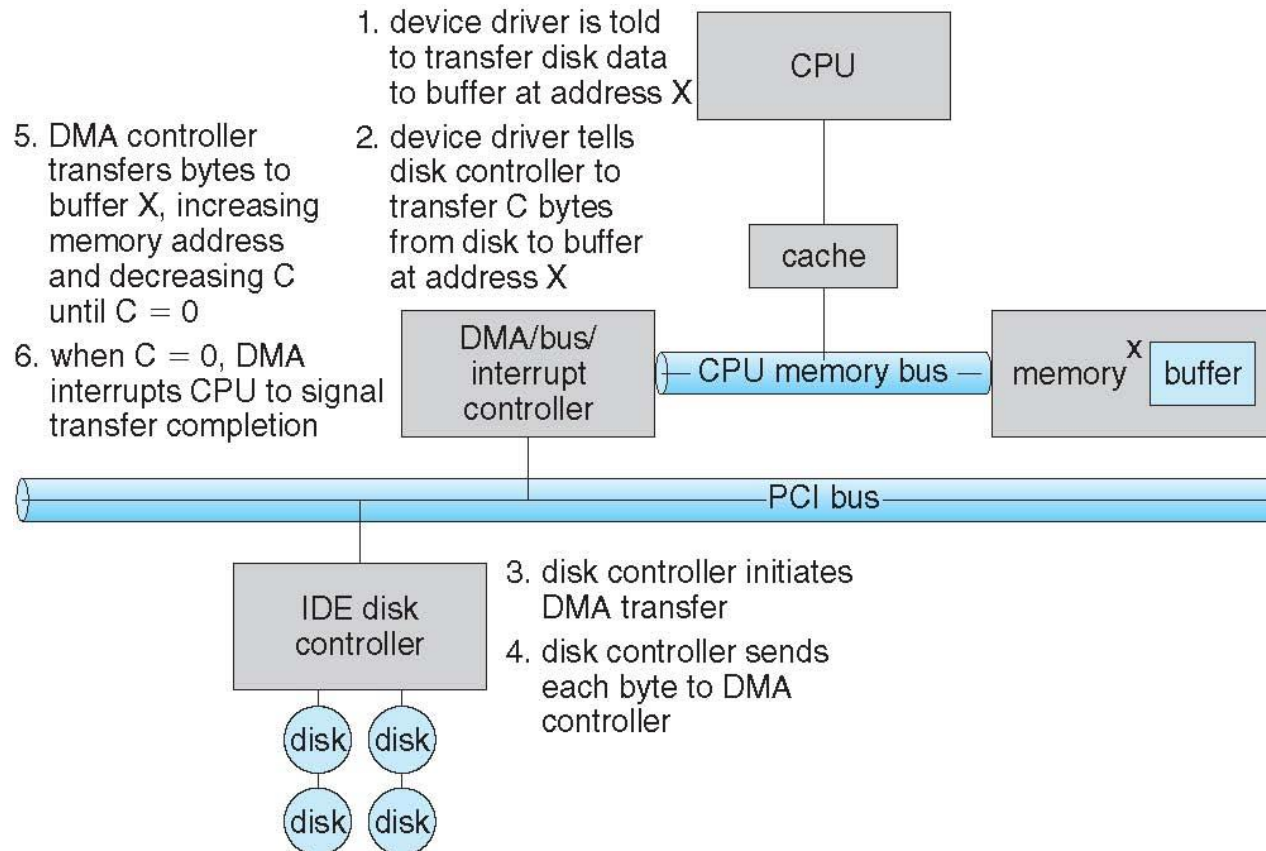
---

- ❁ Handshaking between DMA controller & device controller
  - ❁ Device controller **raises DMA-request** when one word is available
  - ❁ DMA controller **seizes memory bus**, places the desired address on memory-address wires, and raises DMA-acknowledge
  - ❁ Device controller **transfers** the word to memory, and removes the DMA-request signal. Goto 1
  - ❁ DMA controller **interrupts** the CPU.



# Direct Memory Access (DMA方式)

## ❁ Six Step Process to Perform DMA Transfer



❁ Cycle stealing: when DMA seizes the memory bus, CPU is momentarily prevented from accessing main memory





# *Catalog Description*

---

- ✦ I/O Hardware
- ✦ Application I/O Interface
- ✦ Kernel I/O Subsystem
- ✦ Transforming I/O Requests to Hardware Operations
- ✦ Performance



# *Application I/O Interface*

---

- ⊙ Block and Character Devices
- ⊙ Network Devices
- ⊙ Clocks and Timers
- ⊙ Blocking (阻塞) and Nonblocking (非阻塞) I/O



# *I/O control challenges*

---

- Wide variety of devices

- Two challenges

Applications → OS ← Devices

- How can the OS give a convenient, uniform I/O interface to applications?


- How can the OS be designed such that new devices can be attached to the computer without the OS being rewritten?

- For device manufacturers, **device-driver layer** hides differences among I/O controllers from kernel



# *I/O control challenges*

## ❁ A Kernel I/O Structure





# *Application I/O Interface*

- ✿ For applications, I/O system calls encapsulate device behaviors in generic classes
- ✿ 设备独立性：应用程序与具体的物理设备无关。
- ✿ Device-driver layer hides differences among I/O controllers from kernel
- ✿ Devices vary in many dimensions
  - ✦ Character-stream or block
  - ✦ Sequential or random-access
  - ✦ Sharable or dedicated
  - ✦ Speed of operation
  - ✦ read-write, read only, or write only



# *Characteristics of I/O Devices*

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read-write	CD-ROM graphics controller disk



# *Block and Character Devices*

---

- ✿ Block devices include disk drives
  - ✦ Commands include read, write, seek
  - ✦ Raw I/O or file-system access
  - ✦ Memory-mapped file access possible
  
- ✿ Character devices include keyboards, mice, serial ports
  - ✦ Commands include `get()`, `put()`
  - ✦ Libraries layered on top allow line editing



# *Network Devices*

---

- ❁ Varying enough from block and character to have own interface
- ❁ Unix and Windows NT/9x/2000 include socket interface
  - ❖ Separates network protocol from network operation
  - ❖ Server - socket, bind, listen, accept
  - ❖ Client - socket, connect
  - ❖ Includes select() functionality
- ❁ Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)





# *Clocks and Timers*

- ⊙ Provide current time, elapsed time, timer
- ⊙ Hardware clocks
  - ⊠ Real Time Clock (RTC, 实时时钟)
  - ⊠ Time Stamp Counter (TSC, 时间戳计数器)
  - ⊠ Programmable Interval Timer (PIT, 可编程间隔定时器)
    - ✓ used for timings, periodic interrupts
- ⊙ `Ioctl()` (on UNIX) covers odd aspects of I/O such as clocks and timers



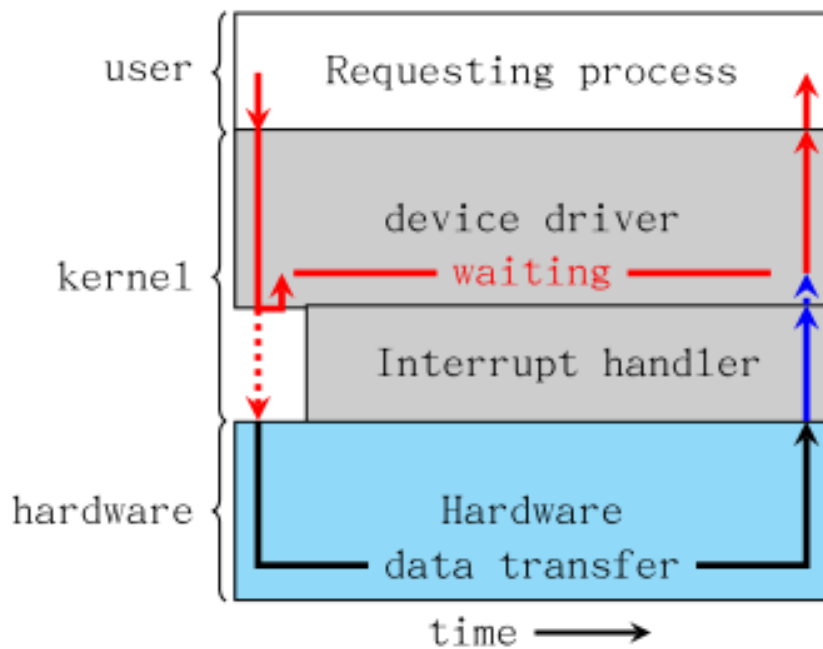
# *Blocking (阻塞) and Nonblocking (非阻塞) I/O*

---


- ⊗ Blocking (阻塞) - suspended until I/O completed
  - ⊗ Easy to use and understand
  - ⊗ Insufficient for some needs
- ⊗ Nonblocking (非阻塞) - I/O call returns as much as
- ⊗ available
  - ⊗ User interface, data copy (buffered I/O)
  - ⊗ Implemented via multi-threading
  - ⊗ Returns quickly with count of bytes read or written
  - ⊗ Asynchronous (异步) - process runs while I/O executes
    - ✓ Difficult to use
    - ✓ I/O subsystem signals process when I/O completed



# Two I/O Methods



(a)  
Synchronous



(b)  
Asynchronous



# *Catalog Description*

---

- ❖ I/O Hardware
- ❖ Application I/O Interface
- ❖ Kernel I/O Subsystem
- ❖ Transforming I/O Requests to Hardware Operations
- ❖ Performance



# *Kernel I/O Subsystem*

---

- ✿ I/O Scheduling
- ✿ Buffering (缓冲机制)
- ✿ Caching, Spooling & device reservation
- ✿ Error Handling
- ✿ I/O Protection
- ✿ Kernel Data Structures



# *I/O scheduling*

---


- ❁ I/O scheduling: To schedule a set of I/O requests means to determine a good order in which to execute them
  - ❁ Origin order: the order in which applications issue system calls: May NOT be the best order!
  - ❁ Scheduling can
    - ✓ Improve overall system performance
    - ✓ Share device access fairly among processes
    - ✓ Reduce the average waiting time for I/O to complete



# *I/O scheduling*

- OS maintaining a wait queue of request for each device

- Device-status Table



- I/O scheduling, Some OSes try fairness, some not



# *I/O scheduling*

---

- ✿ Another way to improve performance is by using storage space in main memory or on disk
  - ✦ Buffering (缓冲机制)
  - ✦ Caching
  - ✦ Spooling





# *Buffering (缓冲机制)*

---

- ⊗ Buffer - A memory area that stores data while they are transferred between two devices or between a device and an application
- ⊗ Store data in memory while transferring between devices
- ⊗ Why buffering?
  - ⊗ To cope with device **speed** mismatch.  
Example: Receive a file via modem and store the file to local hard disk.
    - ✓ Speed: The modem is about a thousand times slower than the hard disk.
    - ✓ Two buffers are used.



# *Buffering (缓冲机制)*

---

## ❁ Why buffering?

- ❁ To cope with device transfer **size** mismatch.

Example: Send/receive a large message via network.

- ✓ At sending side: the large message is fragmented into small network packets.

- ✓ At receiving side: the network packets are placed in a reassembly buffer.

- ❁ To maintain “copy semantics”

Example: When `write()` data to disk, it first copy the data from application's buffer to a kernel buffer.



# *Caching, Spooling & device reservation*


---

- ✿ Caching – fast memory holding copy of data
  - ✦ Always just a copy
  - ✦ Key to performance
- ✿ Spooling – hold output for a device
  - ✦ Dedicated device can serve only one request at a time
  - ✦ Spooling is a way of dealing with I/O devices in a multiprogramming system
  - ✦ Example: Printing
- ✿ Device reservation – provides exclusive access to a device
  - ✦ System calls for allocation and deallocation
  - ✦ Watch out for deadlock



# Spooling

- Out-line I/O (脱机I/O), 使用外围机 (peripheral machine)



- SP00L:  
Simultaneous Peripheral Operation On-Line  
(外部设备联机并行操作, 假脱机)
  - ❑ Dedicated device → sharable device
  - ❑ Using processes of multiprogramming system




# Spooling

## SP00L:

Simultaneous Peripheral Operation On-Line  
(外部设备联机并行操作, 假脱机)

### Structure

- Input-well (输入井), output-well (输出井)
- Input-buffer, output-buffer
- Input-process SP<sub>in</sub>, output-process SP<sub>out</sub>
- Requested-queue





# *Error Handling*

---

- ❊ OS can **recover** from disk read, device unavailable, transient write failures
  - ❏ Example: `read()` again, `resend()`, . . . , according to some specified rules
- ❊ Most return an **error number** or code when I/O request fails
- ❊ **System error logs** hold problem reports




# *I/O Protection*

❖ User process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions

❖ To prevent users from performing illegal I/O

- ❖ All I/O instructions defined to be privileged
- ❖ I/O must be performed via system calls
  - ✓ Memory-mapped and I/O port memory locations must be protected too



Use of a System Call to Perform I/O



# *Kernel Data Structures*

---

- ❁ Kernel **keeps state** info for I/O components, including
  - ❁ open file tables,
  - ❁ network connections,
  - ❁ character device state
- ❁ Many, many complex data structures to **track** buffers, memory allocation, “dirty” blocks
- ❁ Some use object-oriented methods and message passing to implement I/O





# *Catalog Description*

---

- ✦ I/O Hardware
- ✦ Application I/O Interface
- ✦ Kernel I/O Subsystem
- ✦ Transforming I/O Requests to Hardware Operations
- ✦ Performance



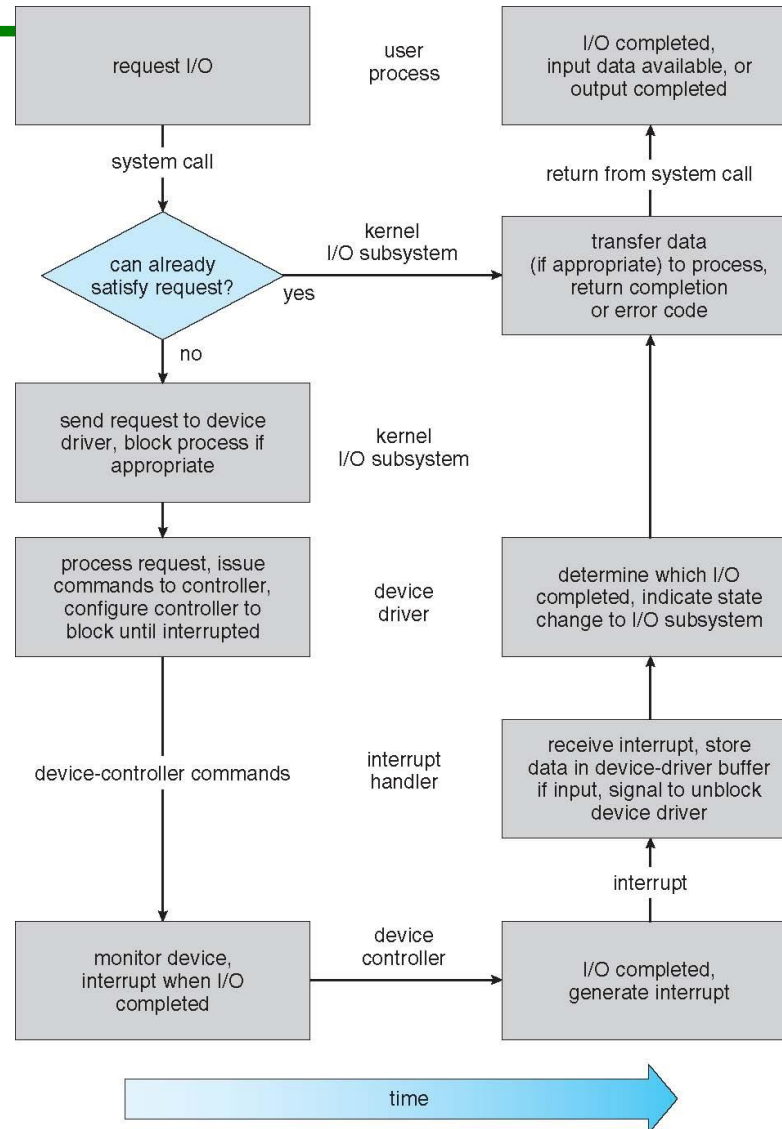
# *I/O Requests to Hardware Operations*

---

- ✿ Consider reading a file from disk for a process:
  - ✦ Determine device holding file
  - ✦ Translate name to device representation
  - ✦ Physically read data from disk into buffer
  - ✦ Make data available to requesting process
  - ✦ Return control to process



# Life Cycle of An I/O Request





# *Catalog Description*

---

- ✦ I/O Hardware
- ✦ Application I/O Interface
- ✦ Kernel I/O Subsystem
- ✦ Transforming I/O Requests to Hardware Operations
- ✦ Performance




# *Performance*

---

- ❁ I/O is a major factor in system performance:
  - ❁ Demands CPU to execute device driver, kernel I/O code
  - ❁ Context switches due to interrupts
  - ❁ Data copying
  - ❁ Network traffic especially stressful



# Intercomputer Communications



Network traffic can also cause a high context-switch rate



# *Improving Performance*

---

- ✿ Reduce number of context switches
- ✿ Reduce data copying
- ✿ Reduce interrupts by using large transfers, smart controllers, polling
- ✿ Use DMA
- ✿ Move processing primitives into hardware
- ✿ Balance CPU, memory, bus, and I/O performance for highest throughput



# *Device-Functionality Progression*

